ASG and LB Session

1.Differences between ELB, ALB, and NLB. Where will you use which one?

“For http protocol require work i choose ALB and for TCP/UDP i choose NLB, ELB is now obselete.”

The Elastic Load Balancer (ELB) was released by AWS in the spring of 2009. An ELB is a software-based load balancer which can be set up and configured in front of a collection of AWS Elastic Compute (EC2) instances. The load balancer serves as a single entry point for consumers of the EC2 instances and distributes incoming traffic across all machines available to receive requests.

In addition to providing a single point of entry, the ELB also performs a vital role in improving the fault tolerance of the services which it fronts. The ELB regularly conducts a health check of all instances which have been registered with it, and only routes traffic to those machines which respond as active and healthy to the health check.

Diff between ALB and NLB:

The first difference is that the Application Load Balancer (as the name implies) works at the Application Layer (Layer 7 of the OSI model). The network load balancer works at layers 3 & 4 (network and transport layers). The network load balancer just forward requests whereas the application load balancer examines the contents of the HTTP request header to determine where to route the request. So, the application load balancer is performing content based routing.

The other difference between the two is important because network load balancing cannot assure availability of the *application.* This is because it bases its decisions solely on network and TCP-layer variables and has no awareness of the application at all. Generally a network load balancer will determine “availability” based on the ability of a server to respond to ICMP ping, or to correctly complete the three-way TCP handshake. An application load balancer goes much deeper, and is capable of determining availability based on not only a successful HTTP GET of a particular page but also the verification that the *content* is as was expected based on the input parameters.

This is also important to note when considering the deployment of multiple applications on the same host sharing IP addresses (virtual hosts in old school speak). A network load balancer will not differentiate between Application A and Application B when checking availability (indeed it cannot unless ports are different) but an application load balancer will differentiate between the two applications by examining the application layer data available to it. This difference means that a

network load balancer may end up sending requests to an application that has crashed or is offline, but an application load balancer will never make that same mistake.

2.Differences between step scaling and target scaling.

(**SIMPLE SCALING**)

You pick ANY Cloud Watch metric

For this and other examples in THIS POST I am choosing CPU Utilization

You specify, a SINGLE THRESHOLD beyond which you want to scale and specify your response

EXAMPLE: how many EC2 instances do you want to add or take away when the CPU UTILIZATION breaches the threshold.

The scaling policy then acts.

THRESHOLD - add 1 instance when CPU Utilization is between 40% and 50%

NOTE: **This is the ONLY Threshold**

(**STEP SCALING**)

You specify MULTIPLE thresholds Along with different responses.

Threshold A - add 1 instance when CPU Utilization is between 40% and 50%

Threshold B - add 2 instances when CPU Utilization is between 50% and 70%

Threshold C - add 3 instances when CPU Utilization is between 70% and 90%

And so on and so on

Note: **There are multiple thresholds**

3.Differences between Launch configuration and launch template.

A launch template is similar to a launch configuration, in that it specifies instance configuration information. Included are the ID of the Amazon Machine Image (AMI), the instance type, a key pair, security groups, and the other parameters that you use to launch EC2 instances. However, defining a launch template instead of a launch configuration allows you to have multiple versions of a template. With versioning, you can create a subset of the full set of parameters and then reuse it to create other templates or template versions. For example, you can create a default template that defines common configuration parameters such as tags or network configurations, and allow the other parameters to be specified as part of another version of the same template.

4.Differences between EC2 healthcheck and load balancer health check

EC2 health check watches for instance availability from hypervisor and networking point of view. For example, in case of a hardware problem, the check will fail. Also, if an instance was misconfigured and doesn't respond to network requests, it will be marked as faulty.

ELB health check verifies that a specified TCP port on an instance is accepting connections OR a specified web page returns 2xx code. Thus ELB health checks are a little bit smarter and verify that actual app works instead of verifying that just an instance works.

That being said there is a third check type: custom health check. If your application can't be checked by simple HTTP request and requires advanced test logic, you can implement a custom check in your code and set instance health though API

5.Create 2 auto-scaling groups with

- launch configuration and

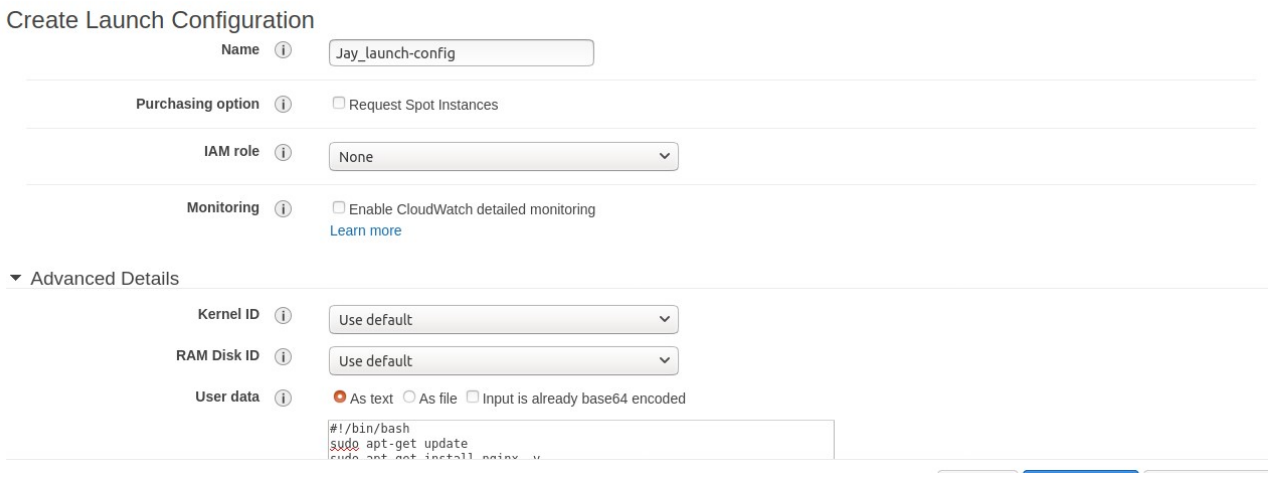

Create Launch Configuration

Review the details of your launch configuration. You can go back to edit the details of each section before you finish.

AMI Details

Edit AMI

Ubuntu Server 18.04 LTS (HVM), SSD Volume Type - ami-07ebfd5b3428b6f4d

Free tier eligible

Ubuntu Server 18.04 LTS (HVM),EBS General Purpose (SSD) Volume Type. Support available from Canonical (http://www.ubuntu.com/cloud/services).

Root device type: ebs Virtualization Type: hvm

Instance Type

Edit instance type

Instance Type	ECUs	vCPUs	Memory GiB	Instance Storage (GiB) GiB	EBS-Optimized Available	Network Performance
t2.micro	Variable	1	1	EBS only	-	Low to Moderate

Launch configuration details

Edit details

Name Jay_launch-config

Cancel | Previous | Create launch configuration

Create Auto Scaling Group

Please review your Auto Scaling group details. You can go back to edit changes for each section. Click **Create Auto Scaling group** to complete the creation of an Auto Scaling group.

Auto Scaling Group Details

Edit details

Group name	Jay-ASG-launch_config
Group size	1
Minimum Group Size	1
Maximum Group Size	1
Subnet(s)	subnet-0303bef94e1bd9734,subnet-0174216d1366316ed
Health Check Grace Period	300
Detailed Monitoring	No
Instance Protection	None
Service-Linked Role	AWSServiceRoleForAutoScaling

Scaling Policies

Edit scaling policies

Notifications

Edit notifications

Tags

Edit tags

Cancel | Previous | Create Auto Scaling group

Auto Scaling group creation status

Successfully created Auto Scaling group

View creation log

View

View your Auto Scaling groups

View your launch configurations

Here are some helpful resources to get you started

Close

Name	Launch Configuration /	Instances	Desired	Min	Max	Availability Zones	Default Cooldown	Health Chec
Jay-ASG-laun…	Jay-launch-config	1	1	1	1	us-east-1c	300	300

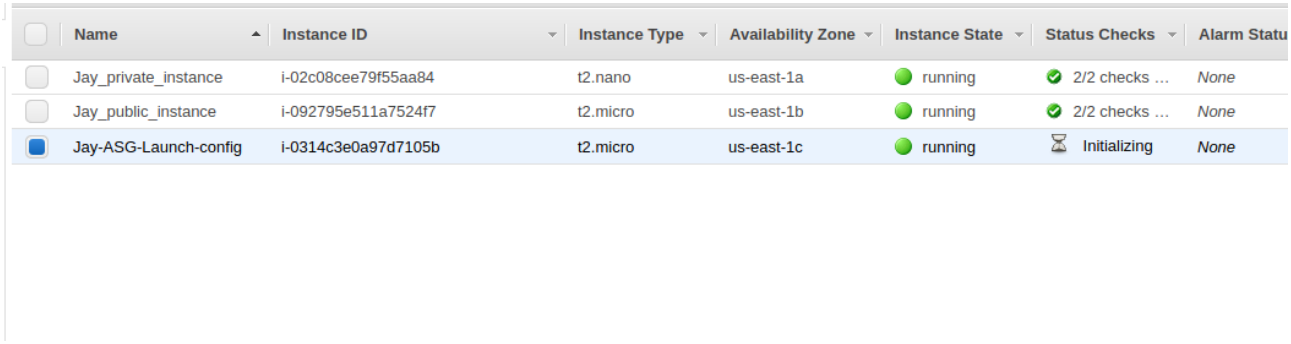

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Statu
Jay_private_instance	i-02c08cee79f55aa84	t2.nano	us-east-1a	running	2/2 checks ...	None
Jay_public_instance	i-092795e511a7524f7	t2.micro	us-east-1b	running	2/2 checks ...	None
Jay-ASG-Launch-config	i-0314c3e0a97d7105b	t2.micro	us-east-1c	running	Initializing	None

- launch template

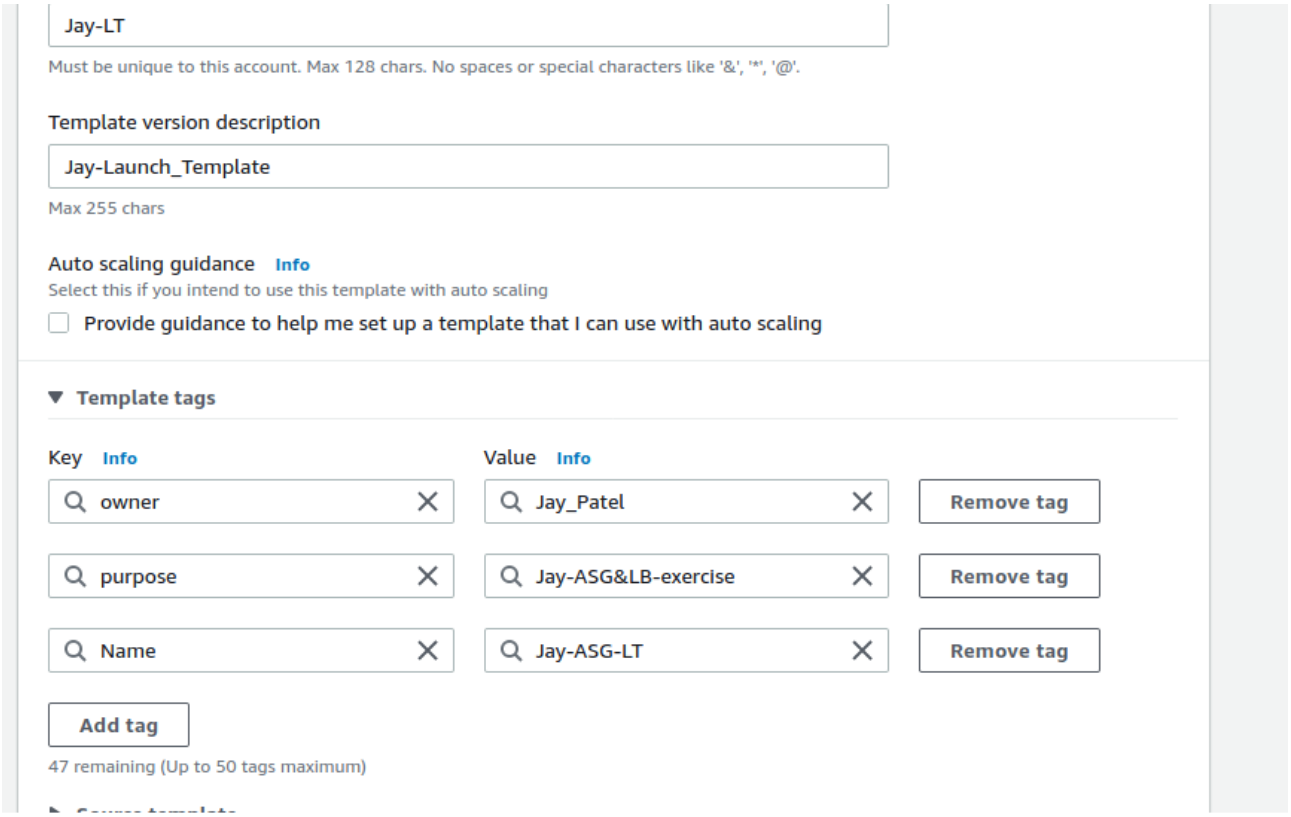

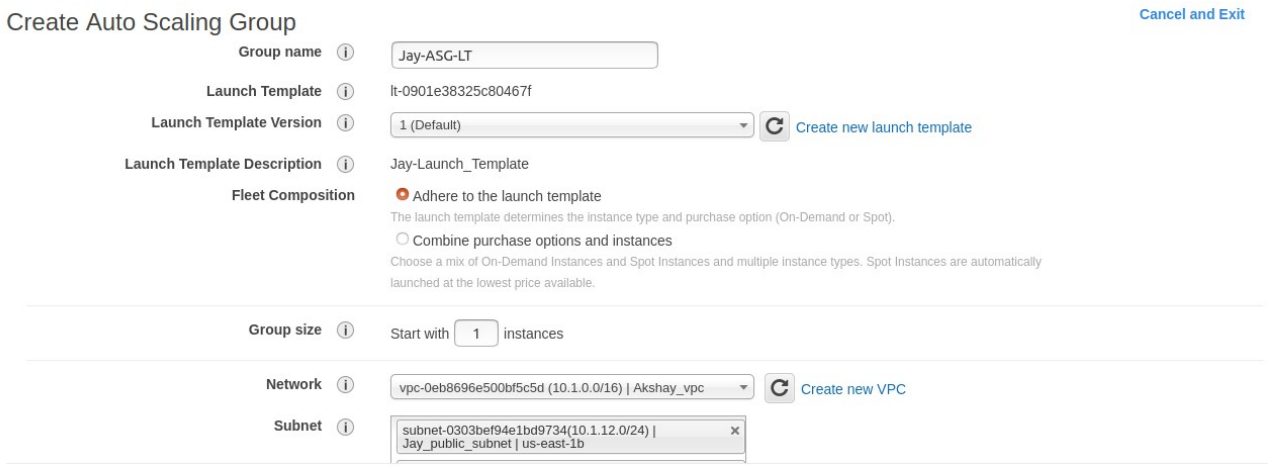

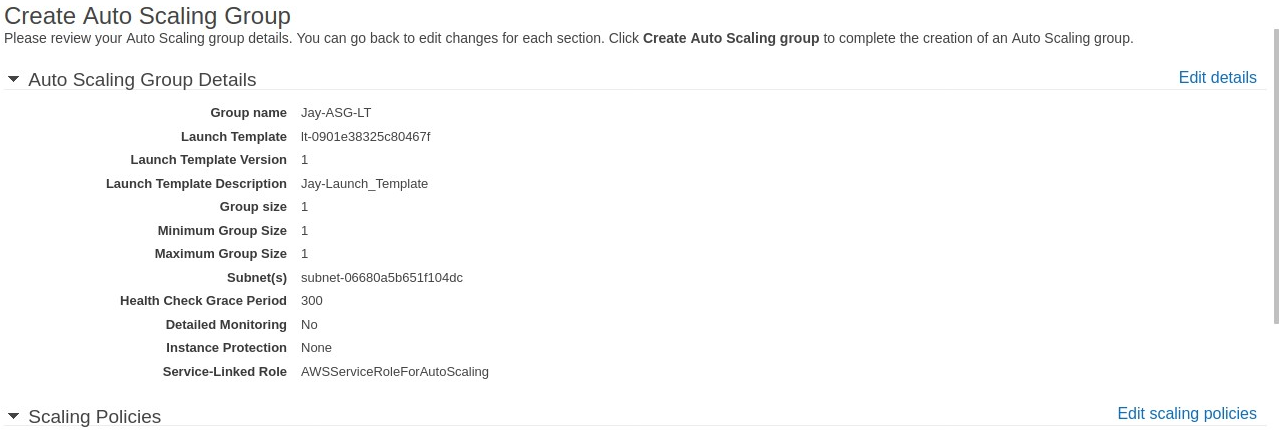

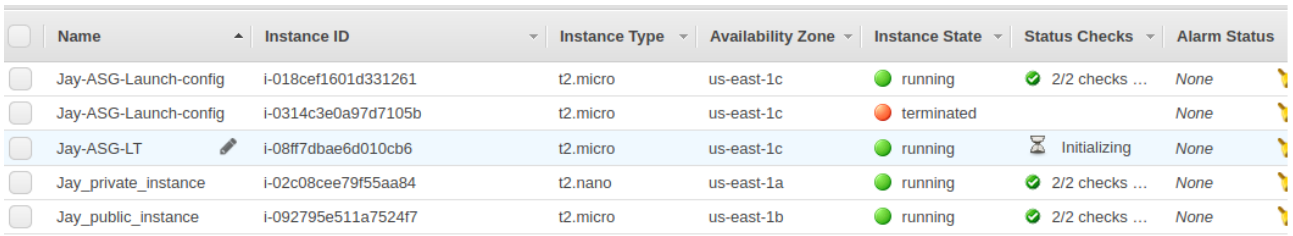

6.Setup autoscaling Wordpress application with the Application load balancer. Auto-scaling should be triggered based on CPU usage of EC2 instances.

Step 1: Configure Load Balancer

Basic Configuration

To configure your load balancer, provide a name, select a scheme, specify one or more listeners, and select a network. The default configuration is an Internet-facing load balancer in the selected network with a listener that receives HTTP traffic on port 80.

Name Jay-ASG-ALB

Scheme (•) internet-facing ( ) internal

IP address type ipv4

Listeners

A listener is a process that checks for connection requests, using the protocol and port that you configured.

Load Balancer Protocol	Load Balancer Port
HTTP	80

Cancel | Next: Configure Security Settings

Step 6: Review

Load balancer — Edit

Name Jay-ASG-ALB
Scheme internet-facing
Listeners Port:80 - Protocol:HTTP
IP address type ipv4
VPC vpc-0eb8696e500bf5c5d (Akshay_vpc)
Subnets subnet-0303bef94e1bd9734 (Jay_public_subnet), subnet-0174216d1366316ed (aks-pub1)
Tags onwer:Jay_Patel, purpose:Jay-ASG&LB-exercise, Name:Jay-ASG-ALB

Security groups — Edit

Security groups Jay-ASG-ALB

Routing — Edit

Target group New target group
Target group name wordpress
Port 80
Target type instance
Protocol HTTP
Health check protocol HTTP

Create Auto Scaling Group

Please review your Auto Scaling group details. You can go back to edit changes for each section. Click **Create Auto Scaling group** to complete the creation of an Auto Scaling group.

Auto Scaling Group Details — Edit details

Group name Jay-ASG-ALB
Launch Template lt-0901e38325c80467f
Launch Template Version 2
Launch Template Description Jay-LT-wordpress
Group size 1
Minimum Group Size 1
Maximum Group Size 1
Subnet(s) subnet-0303bef94e1bd9734,subnet-0174216d1366316ed
Health Check Grace Period 300
Detailed Monitoring No
Instance Protection None
Service-Linked Role AWSServiceRoleForAutoScaling

Scaling Policies — Edit scaling policies

Cancel | Previous | Create Auto Scaling group

Name	Launch Configuration /	Instances	Desired	Min	Max	Availability Zones	Default Cooldown	Health Chec
Jay-ASG-ALB	Jay-LT	1	1	1	3	us-east-1b, us-east-1d	300	300
Revant_wordp…	Revant_wordpress	1	1	1	2	us-east-1c	300	300
Vedant-test-gr…	Vedant-Test-Config	3	3	3	3	us-east-1b, us-east-1c, us-e…	300	300
Test-AGS	Test_Template	2	2	1	5	us-east-1c	300	300

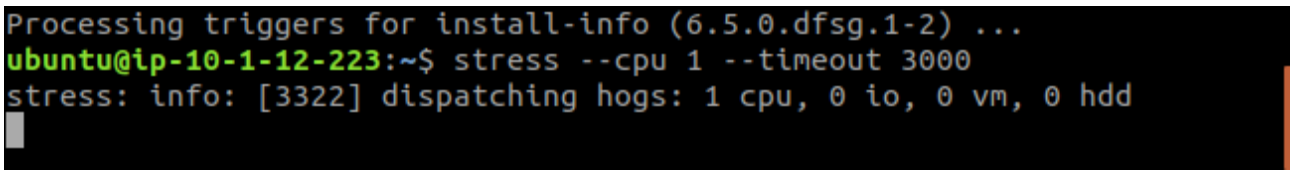

Jay-ASG-ALB	i-093f6c5a3902ef11b	t2.micro	us-east-1b	running	2/2 checks ...	None
Jay-ASG-ALB	i-0f499eece0f39f953	t2.micro	us-east-1d	running	2/2 checks ...	None
Jay-ASG-LT	i-08ff7dbae6d010cb6	t2.micro	us-east-1c	terminated		None
Jay_private_instance	i-02c08cee79f55aa84	t2.nano	us-east-1a	stopped		None
Jay_public_instance	i-092795e511a7524f7	t2.micro	us-east-1b	running	2/2 checks ...	None

Name	Launch Configuration /	Instances	Desired	Min	Max	Availability Zones	Default Cooldown	Health
Jay-ASG-ALB	Jay-LT	2	2	1	3	us-east-1b, us-east-1d	300	300
Vedant-test-gr...	Vedant-Test-Config	3	3	3	5	us-east-1b, us-east-1c, us-e...	300	300
revant	revant	1	1	1	2	us-east-1b, us-east-1e	300	300
Test-AGS	Test_Template	2	2	1	5	us-east-1c	300	300
abhishek	Abhishek	0	0	0	0	us-east-1a, us-east-1b	300	300

Status	Description	Start Time	End Time
Successful	Launching a new EC2 instance: i-07e81ae7a03a8bb47	2020 February 27 12:16:57 UTC+5:30	2020 February 27 12:17:30 UTC+5:30

Description: Launching a new EC2 instance: i-07e81ae7a03a8bb47

Cause: At 2020-02-27T06:46:24Z a monitor alarm awsec2-Jay-ASG-ALB-CPU-Utilization in state ALARM triggered policy CPUHIGH changing the desired capacity from 1 to 2. At 2020-02-27T06:46:55Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 2.

Status	Description	Cause
Successful	Terminating EC2 instance: i-07e81ae7a03a8bb47	At 2020-02-27T07:01:25Z a monitor alarm awsec2-Jay-ASG-ALB-High-CPU-Utilization i triggered policy CPULow changing the desired capacity from 2 to 1. At 2020-02-27T07:0 instance was taken out of service in response to a difference between desired and actual shrinking the capacity from 2 to 1. At 2020-02-27T07:01:29Z instance i-07e81ae7a03a8 selected for termination.
Successful	Terminating EC2 instance: i-0f499eece0f39f953	At 2020-02-27T06:55:35Z a user request executed policy CPULow changing the desired to 2. At 2020-02-27T06:55:40Z an instance was taken out of service in response to a diff between desired and actual capacity, shrinking the capacity from 3 to 2. At 2020-02-27T instance i-0f499eece0f39f953 was selected for termination.
Successful	Launching a new EC2 instance: i-0b6e029695c626837	At 2020-02-27T06:53:24Z a monitor alarm awsec2-Jay-ASG-ALB-CPU-Utilization in stat triggered policy CPUHIGH changing the desired capacity from 2 to 3. At 2020-02-27T06 instance was started in response to a difference between desired and actual capacity, inc capacity from 2 to 3.
Successful	Launching a new EC2 instance: i-07e81ae7a03a8bb47	At 2020-02-27T06:46:24Z a monitor alarm awsec2-Jay-ASG-ALB-CPU-Utilization in stat triggered policy CPUHIGH changing the desired capacity from 1 to 2. At 2020-02-27T06 instance was started in response to a difference between desired and actual capacity, inc capacity from 1 to 2.
Successful	Terminating EC2 instance: i-093f6c5a3902ef11b	At 2020-02-27T06:39:39Z a user request update of AutoScalingGroup constraints to mir desired: 1 changing the desired capacity from 2 to 1. At 2020-02-27T06:40:37Z an instar out of service in response to a difference between desired and actual capacity, shrinking from 2 to 1. At 2020-02-27T06:40:37Z instance i-093f6c5a3902ef11b was selected for t

7.Create another Wordpress website and use the ALB created above to send traffic to this website based on the hostname

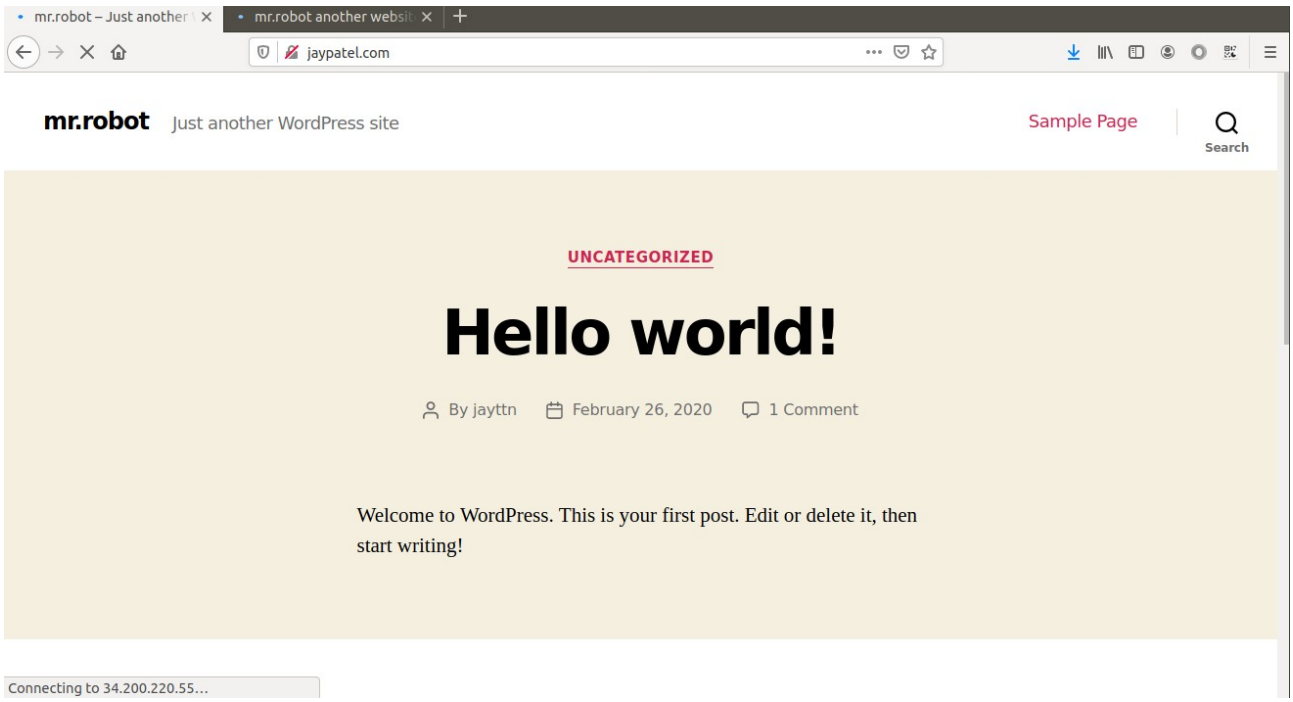
mr.robot – Just another
mr.robot another websit
jaypatel.com
mr.robot
Just another WordPress site
Sample Page
Search
UNCATEGORIZED
Hello world!
By jayttn
February 26, 2020
1 Comment
Welcome to WordPress. This is your first post. Edit or delete it, then start writing!
Connecting to 34.200.220.55...

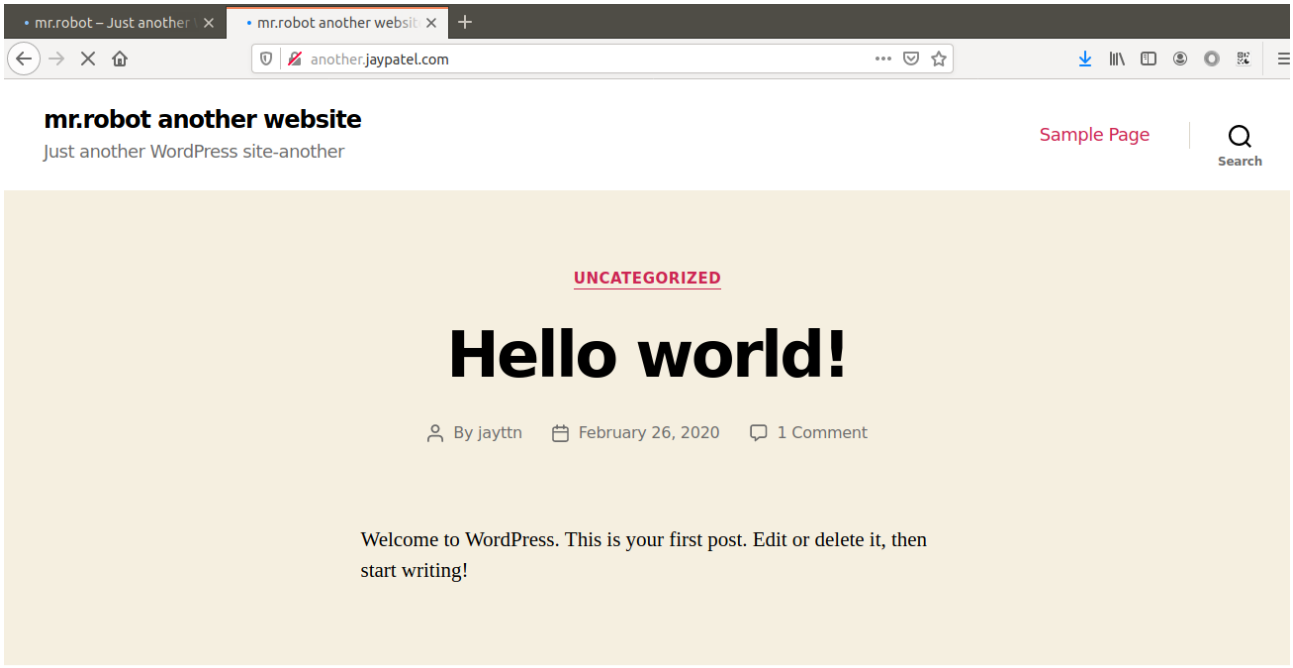
mr.robot – Just another
mr.robot another websit
another.jaypatel.com
mr.robot another website
Just another WordPress site-another
Sample Page
Search
UNCATEGORIZED
Hello world!
By jayttn
February 26, 2020
1 Comment
Welcome to WordPress. This is your first post. Edit or delete it, then start writing!

Jay-ASG-ALB | HTTP:80 (3 rules)

Rule limits for condition values, wildcards, and total rules.

		IF	THEN
1	arn...24661	Host is jaypatel.com	Forward to Jay-wordpress: 1 (100%) Group-level stickiness: Off
2	arn...c3f86	Host is another.jaypatel.com	Forward to Jay-another-wordpress: 1 (100%) Group-level stickiness: Off
last	HTTP 80: default action *This rule cannot*	Requests otherwise not routed	Forward to Jay-wordpress: 1 (100%)

Name	Port	Protocol	Target type	Load Balan	VPC ID	Monitoring
Jay-another-wordpress	80	HTTP	instance	Jay-ASG-A...	vpc-0eb8696e500bf5c5d	
Jay-wordpress	80	HTTP	instance	Jay-ASG-A...	vpc-0eb8696e500bf5c5d	

Target group: Jay-another-wordpress

Description | Targets | Health checks | Monitoring | Tags

Basic Configuration

Name: Jay-another-wordpress

ARN:
arn:aws:elasticloadbalancing:us-east-1:187632318301:targetgroup/Jay-another-wordpress/bdd4ce9e90e446e8

Protocol: HTTP

Port: 80

The load balancer starts routing requests to a newly registered target as soon as the registration process completes and the target passes the initial health checks. If demand on your targets increases, you can register additional targets. If demand on your targets decreases, you can deregister targets.

Edit

Registered targets

Instance ID	Name	Port	Availability Zone	Status	Description
i-092795e511a7524f7	Jay_public_instance	80	us-east-1b	healthy	This target is currently passing target group's health checks.

Availability Zones

Availability Zone	Target count	Healthy?
us-east-1b	1	Yes

Target group: Jay-wordpress

Description | Targets | Health checks | Monitoring | Tags

Edit

Registered targets

Instance ID	Name	Port	Availability Zone	Status	Description
i-00173f2ca747883aa	Jay-ASG-ALB	80	us-east-1b	healthy	This target is currently passing target group's health checks.

Availability Zones

Availability Zone	Target count	Healthy?

8.Use NLB that replaces the ALB in the above setup.

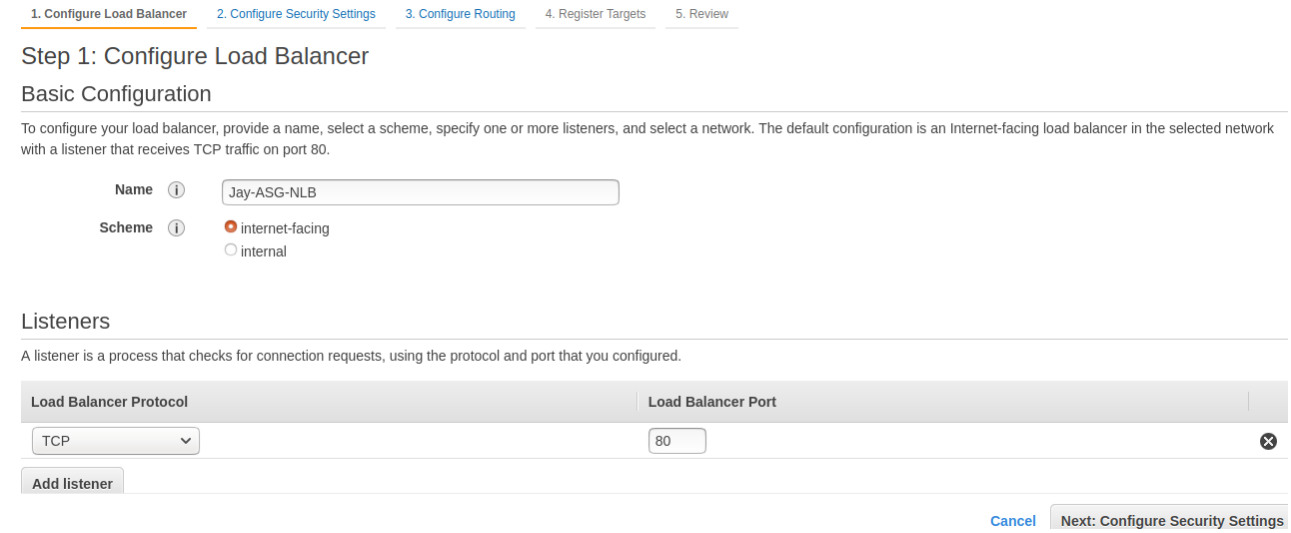

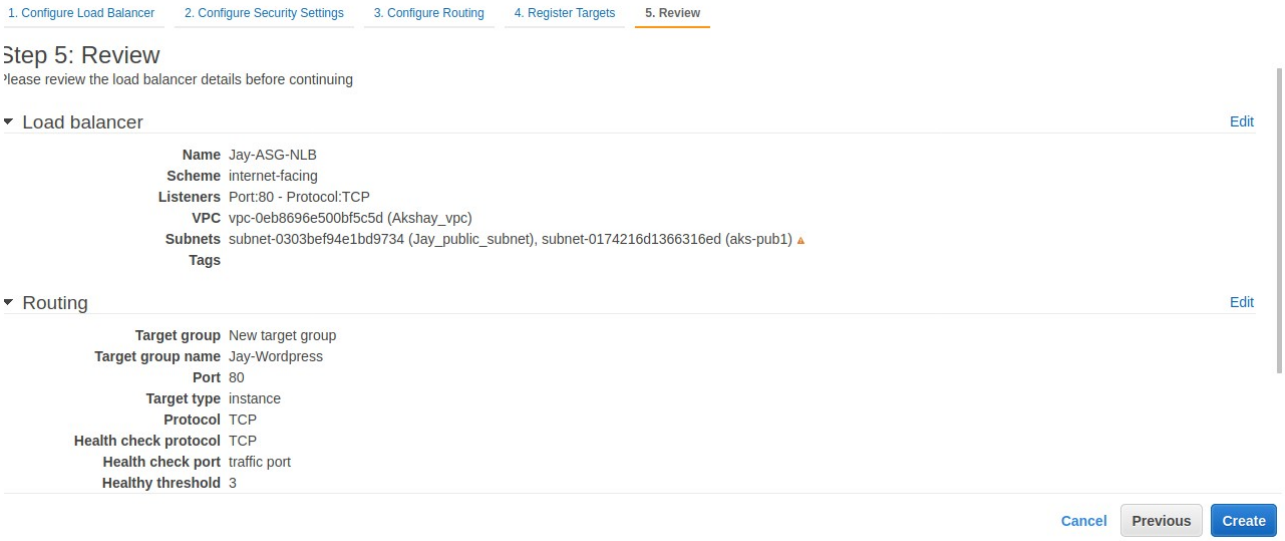

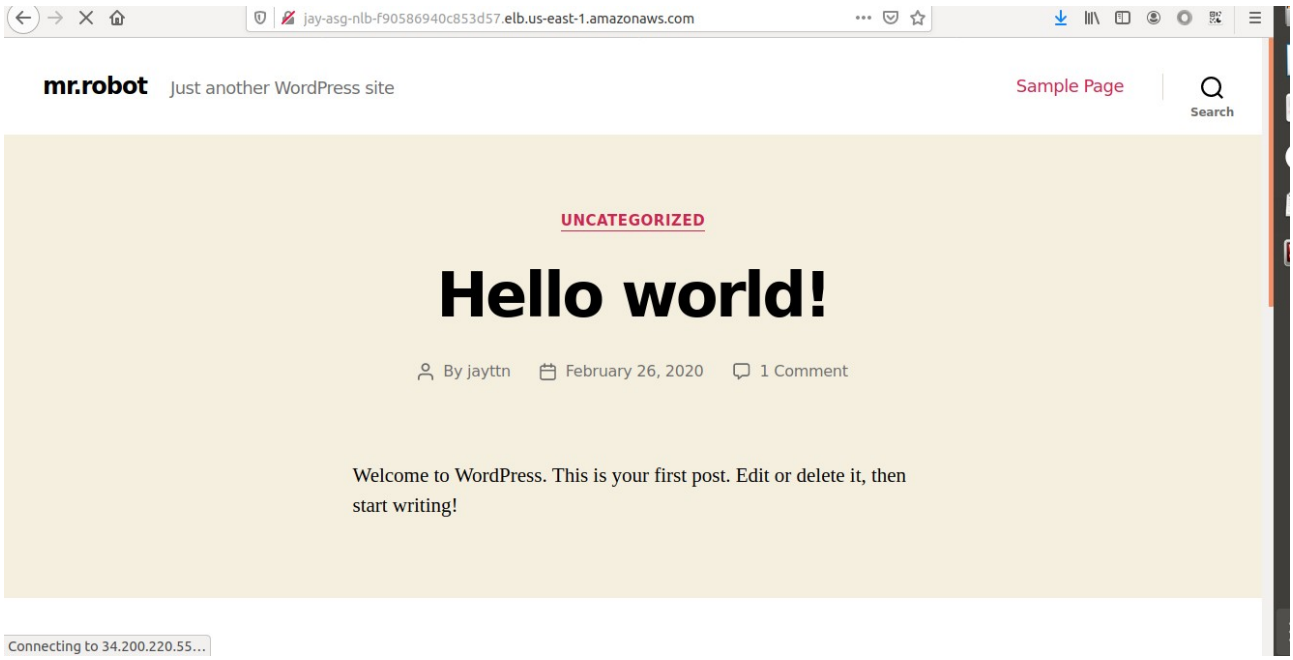

9.Take an instance out of the ASG.

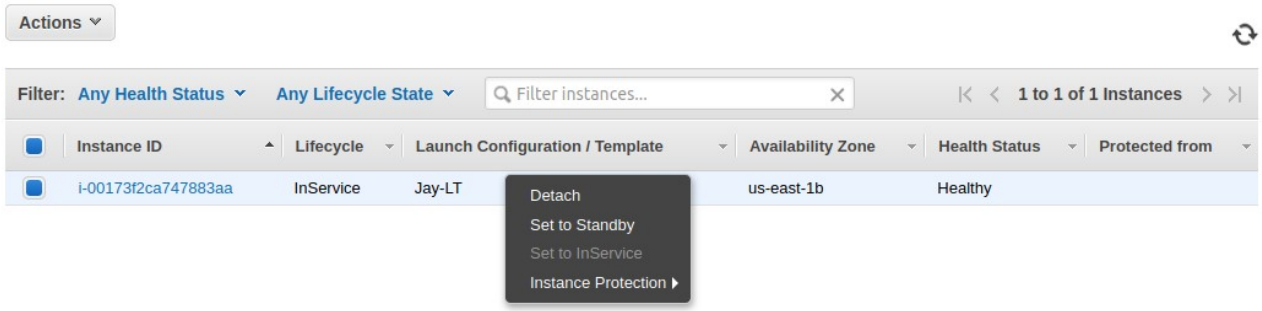

10.Put scale-in protection on an instance in the ASG.

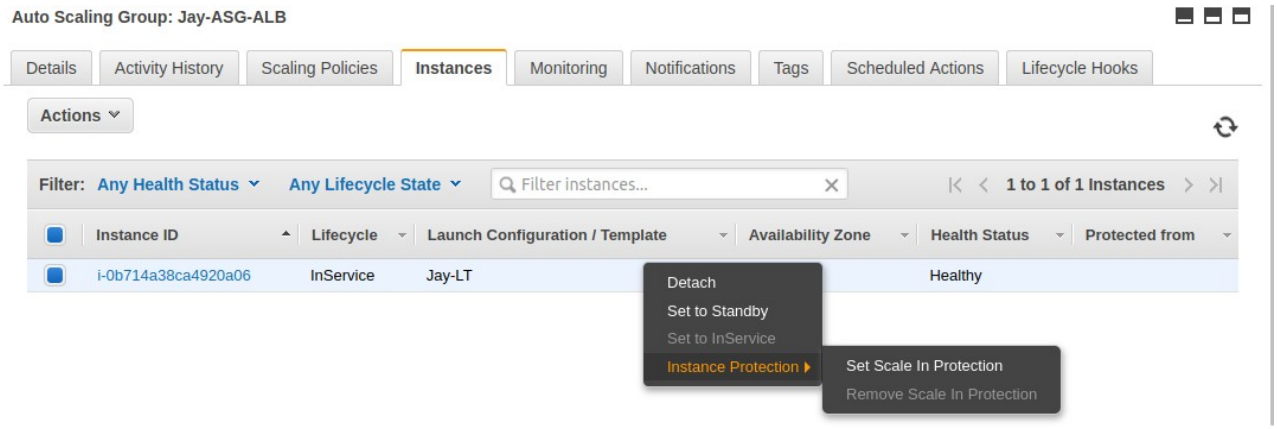

11.Put Schedules in ASG to:

- Remove all instances of the ASG at 8 PM

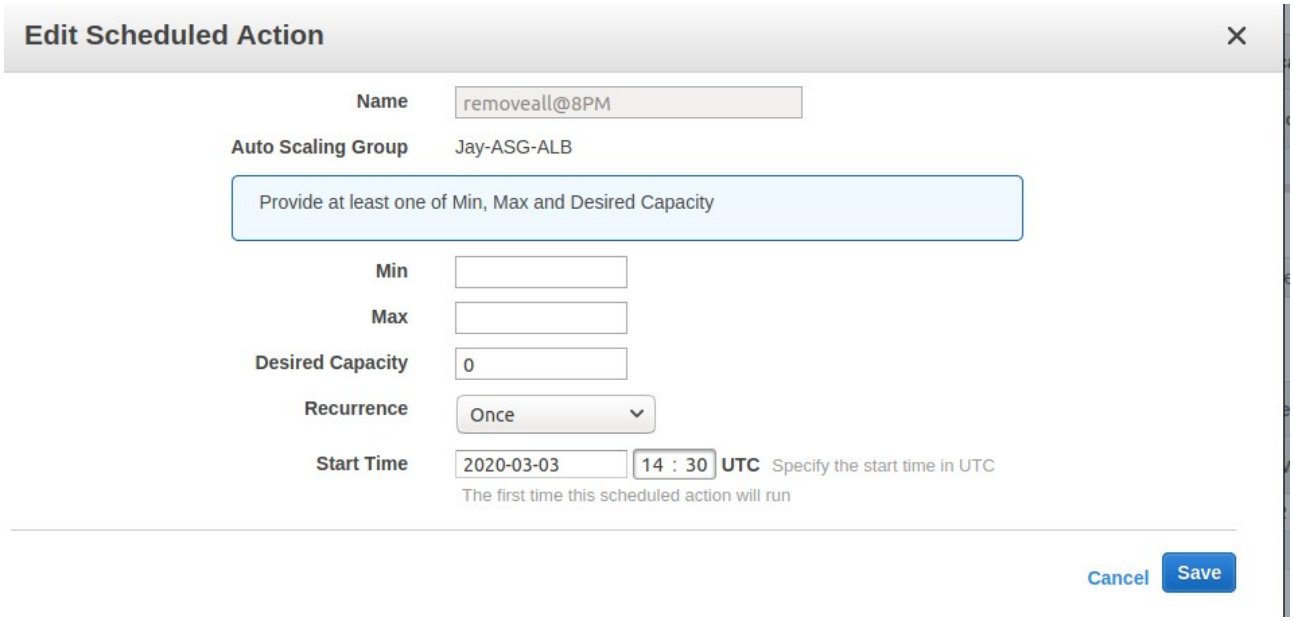

- Launch a minimum of 2 instances at 10 AM

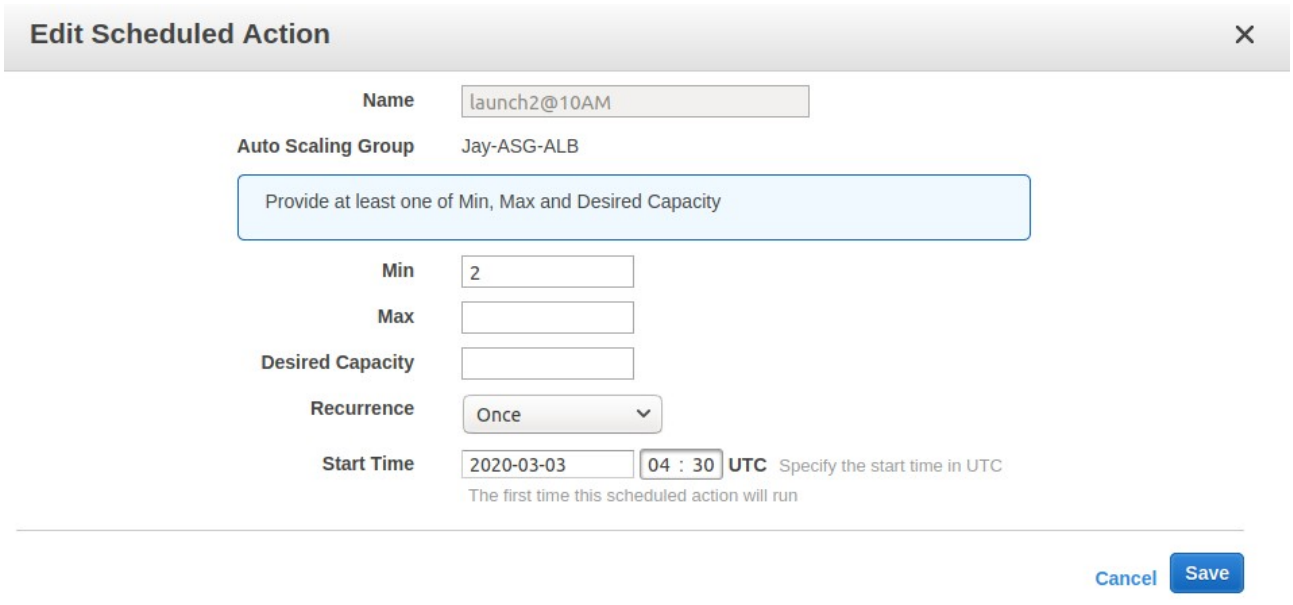

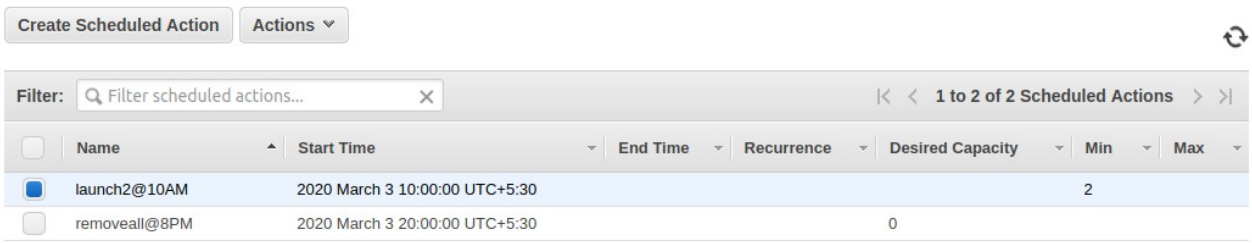
Create Scheduled Action
Actions
Filter:
Filter scheduled actions...
1 to 2 of 2 Scheduled Actions
Name
Start Time
End Time
Recurrence
Desired Capacity
Min
Max
launch2@10AM
2020 March 3 10:00:00 UTC+5:30
2
removeall@8PM
2020 March 3 20:00:00 UTC+5:30
0